

Highlights of the NeurIPS 2023 unlearning challenge


5,161 registrations

1,338 participants from 72 countries.

For 500 (including 44 in the top 100!), this was their first competition

1,121 teams

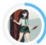
















1,923 submissions



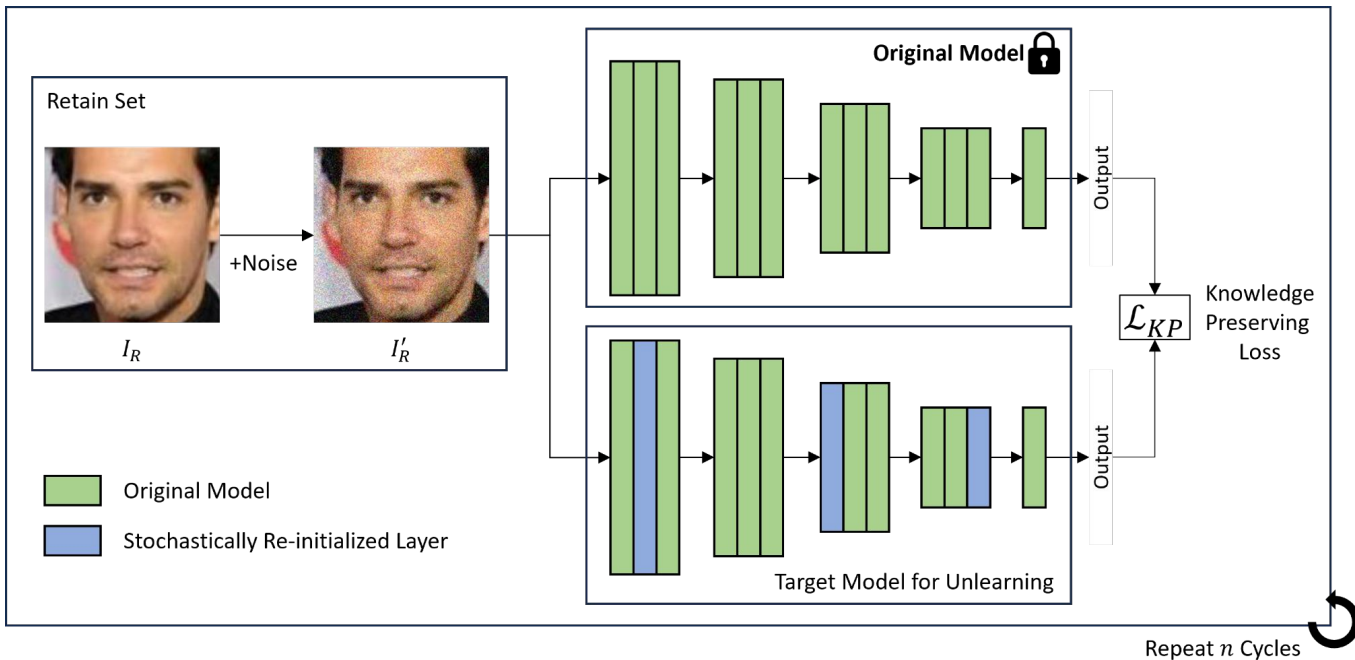
Leaderboard: 40% scored above baseline

Top submissions

■ Prize Winners

#	Team	Members
1	fanchuan	
2	[kookmin Univ] LD&BGW&KJH	 
3	Seif Eddine Achour	
4	Sebastian Oleszko	
5	toshi_k & marvelworld	 
6	Algorithmic Amnesiacs	   
7	Jiaxi Sun	
8	Forget	    

8th place solution - Team Forget



(1) Forgetting phase: model parameters are stochastically selected and re-initialized.

- FC, Projection-shortcut layers are excluded from the selection pool.

(2) Remembering phase: knowledge preserving loss is calculated between the original model and the target unlearning model.

- Knowledge Preserving Loss:

$$\mathbb{E}\{|f_O(I'_R) - f_U(I'_R)|^2\}$$



- Gaussian noise is added to the image as data augmentation.

- It reminds the target model about the retain set.

(3) Forgetting phase and Remembering phase are repeated for n cycles to enhance unlearning performance.

8th place solution - Team Forget

Forget Set





- Histograms of logits from retrained model and unlearned models are visualized.

- This observation is acquired from local experiments on CIFAR-10.

- MSE loss makes closer distributions than CE loss for both forget set and retain set.

Retain Set



(a) CE

(b) MSE

Figure. Comparison of logit distributions between CE loss and MSE loss

8th place solution - Team Forget

Table 1. Comparison between different data augmentation techniques

Input Data	Score↑
Clean Image	0.06172
Vertical Flip	0.02505
Random Crop	0.00001
Cutout	0.00001
Image + Gaussian Noise ($\sigma = 0.1$)	0.06532

- Gaussian Noise ($\sigma=0.1$) is the best data augmentation compared with other data augmentation techniques.

- Compared with CE loss and L1 loss, MSE loss demonstrates the best score.

- Additionally, increasing the cycles highly improves the performance.

- From these observations, we build the final submission.

Table 2. Comparison between different sigma of gaussian noise

Input Data	Score↑
Image + Gaussian Noise ($\sigma = 0.05$)	0.06333
Image + Gaussian Noise ($\sigma = 0.15$)	0.05907
Image + Gaussian Noise ($\sigma = 0.1$)	0.06532

Table 4. Effect of cycles

Number of Cycles	Selection Ratio	Score↑
1 (2 epochs)	10%	0.0680
1 (2 epochs)	20%	0.0656
2 (2-2 epochs)	~20%	0.0844
3 (1-2-2 epochs)	~30%	0.0856

Table 3. Comparison between different loss functions

Loss Function	Score↑
CE Loss	0.0653
L1 Loss	0.0326
MSE Loss	0.0680

Table 5. Final submission score compared with other unlearning methods

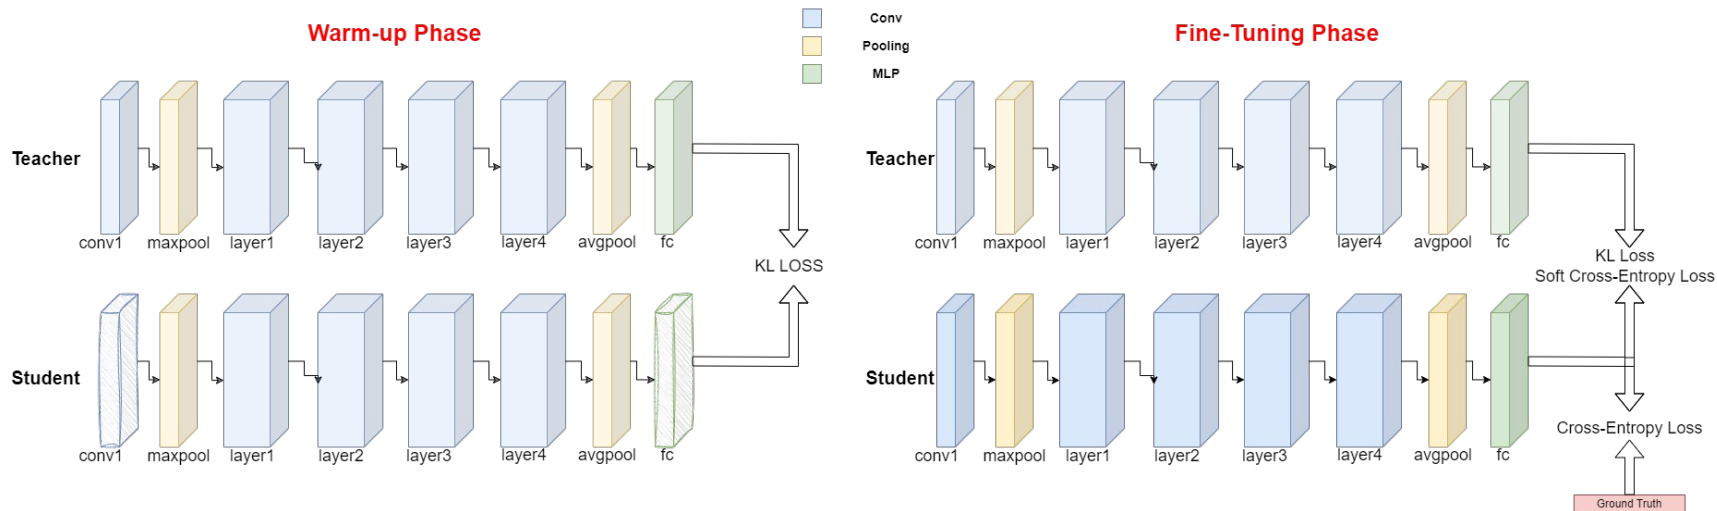
Model	Score↑
Negrad	0.0001 (± 0.0001)
Fine-tune	0.0464 (± 0.0031)
Ours	0.0935 (± 0.0060)
Ours, best	0.1024

7th place solution - Jiaxi Sun

Solution that only makes use of retain set

1. Reset parameters of last layer
1. Randomly selecting $N=9$ layers from the network and add noise
 - a. Adding noise helps the network 'forget' the information it has learned, and the randomness of the layer selection contributes to enhancing the model's diversity.
1. Fine-tune all network layers

6th place solution - Algorithmic Amnesiacs

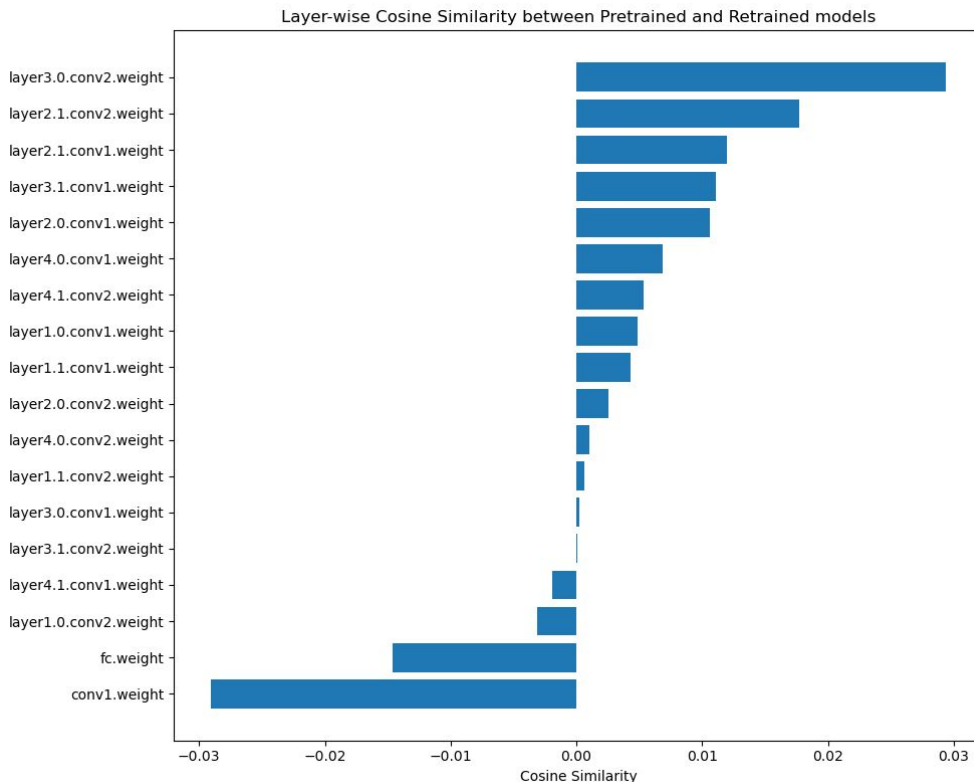


1. Reset first and last layer of the original model.
2. Warm-up phase employing knowledge distillation
3. Fine-tuning phase.

6th place solution - Algorithmic Amnesiacs

Reset first and last layers:


- First layer significantly influences the rest of the model layers and the last layer determines the model's final output distribution.
- On CIFAR-10: these two layers exhibited the most negative cosine similarity between model weights trained on the full training set and models trained from scratch on a smaller subset (i.e., the retain set).



6th place solution - Algorithmic Amnesiacs

Warm-up phase

Minimize KL divergence between the outputs of the original pre-trained model (teacher) and the reinitialized model (student) **on the validation set.**




6th place solution - Algorithmic Amnesiacs

Fine-tuning using 3 losses

Cross-entropy for model's accuracy using hard labels on retain set.

Soft cross-entropy for predictions of the student model with soft labels from the teacher model.

KL divergence combined with the soft cross-entropy facilitates rapid knowledge transfer and broader information capture.



5th place solution - toshi_k & marvelworld

Summary

Our solution is the ensemble of two approaches:

- (1) Retraining from transposed weights
- (2) Fine-tuning with pseudo-labels.


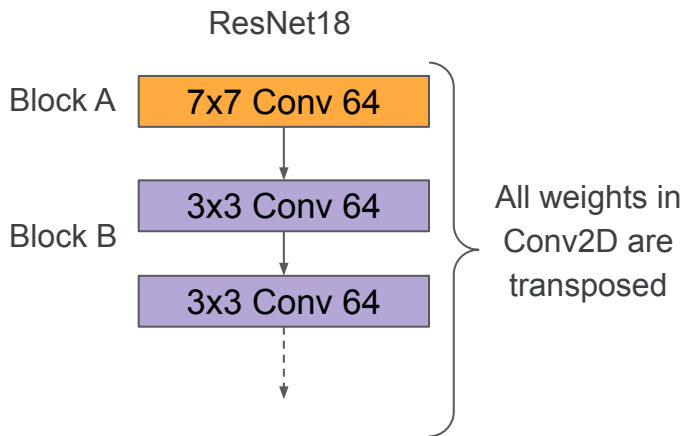
Our solution is built upon two distinctive approaches, contributing to the stability of our solution in the private LB.

(1) Retraining from transposed weights	(2) Fine-tune with pseudo-labels	Public LB	Private LB
512 models	0 models	0.0720386947	-
0 models	512 models	0.0707241647	-
246 models	266 models	-	0.0785184178
266 models	246 models	-	0.0756313425

5th place solution - toshi_k & marvelworld

(1) Retraining from transposed weights

- This part retrains the model using a modified version of the original model.
- In this modification, all weights in Conv2D are transposed. This process helps in forgetting samples in the forget-set, enabling the reuse of valuable features from the original model.



```
for module in local_model.modules():  
    if isinstance(module, torch.nn.modules.conv.Conv2d):  
        module.weight = torch.nn.Parameter(module.weight.swapaxes(2, 3))
```

The modification is carried out simply like this.

5th place solution - toshi_k & marvelworld

(2) Fine-tuning with pseudo-labels

- This part reproduces behavior (errors) on the forget data with pseudo-labels from two functions.

Algorithm 1 Unlearning with pseudo-labels

Input: forget data : $(x_f, y_f) \in D_F$
retrain data : $(x_r, y_r) \in D_R$
pretrained model : f_θ^p
simple finetuning model : $f_\theta^l = f_\theta^p(x_r) \rightarrow \hat{y}^{pr} \rightarrow \text{Loss}(\hat{y}^{pr}, y^r)$
simple scratch model : $f_\theta^s = f_\theta(x_r) \rightarrow \hat{y}^r \rightarrow \text{Loss}(\hat{y}^r, y^r)$
confidence threshold : T

```
function INCORRECT DIRECTION PSEUDO-LABELS
  for  $x_f \in D_F$  do
    if  $f_\theta^p(x_f) = y^f$  and  $f_\theta^p(x_f) \neq f_\theta^l(x_f)$  then
       $D_P \ni \{x_f, \hat{y}^{lf}\}$ 
    end if
  end for
  return  $D_P$ 
end function
```

```
function HIGH CONFIDENCE INCORRECT PSEUDO-LABELS
  for  $x_f \in D_F$  do
    if  $\text{Entropy}(f_\theta^s(x_f)) < T$  and  $f_\theta^s(x_f) \neq y^f$  then
       $D_P \ni \{x_f, \hat{y}^{sf}\}$ 
    end if
  end for
  return  $D_P$ 
end function
```

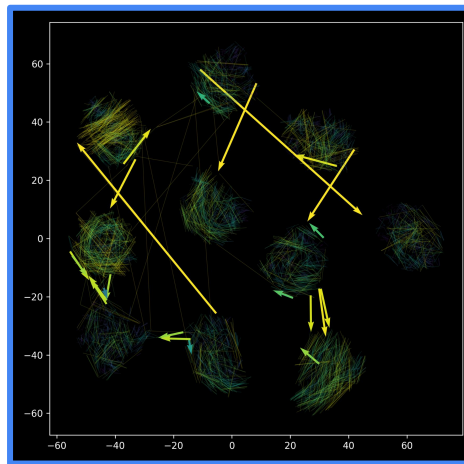


Figure 1: Inference shifts between the pretrained model and the fine-tuned model, showcasing significant shifts in the incorrect direction.

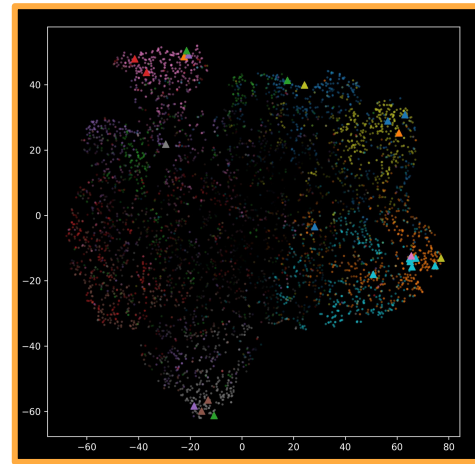


Figure 2: Inference results of the scratch model. Triangle-up marker indicates high confidence but they are making incorrect inferences.

4th place solution - Sebastian Oleszko

1. Re-initializing/pruning 99% of parameters based on L1-norm (Unstructured)
 - Weights: Pytorch default initialization
 - Biases: Set to zero (prune)
2. Fine-tune on retain dataset
 - Regularize using entropy
 - Cross entropy class weights as $N_c^{-0.1}$

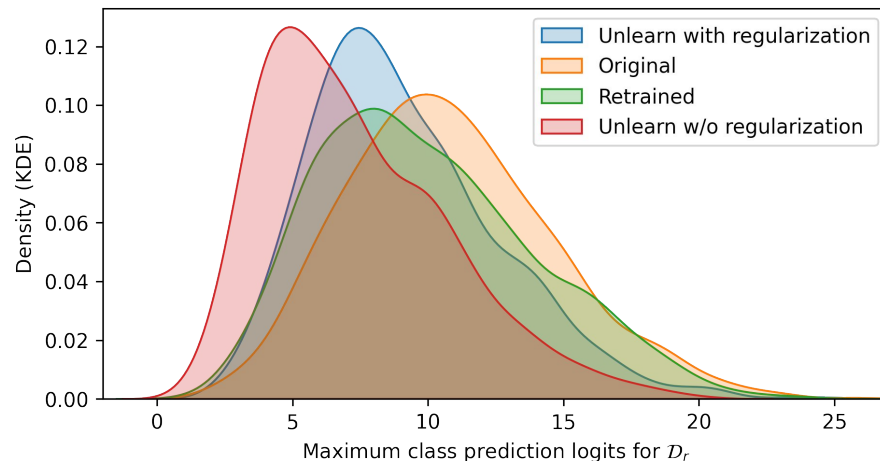

$$\min_w \sum_{(x_r, y_r) \in \mathcal{D}_r} H(y_r, f(x_r; w)) + \sum_{x_r \in \mathcal{D}_r} (H(f(x_r; w)) - H(f(x_r; w^o)))^2$$

Cross-entropy MSE of entropy

Initial weights
↓

4th place solution - Sebastian Oleszko

- CIFAR-10 experiment
- Impact of most important hyperparameters: Learning rate/epochs and pruning percentage
- Effect of including entropy regularization
- Tuning on public submission scores



4th place solution - Sebastian Oleszko

Entropy-based regularization

Helps to achieve a more similar prediction distribution/confidence.

Unlearning through pruning/re-initialization

Effective as unlearning technique. Most of the performance is retainable even with high pruning percentage.


Concluding thoughts

- Hyperparameter tuning is **very** important to achieve high scores
- Final submission was only fine-tuned for 3.2 epochs - maybe not optimal

3rd place solution - Seif Eddine Achour


Competition approach: vision confusion - reconstitution

Vision Confusion




Details forgotten but
the general idea was
retained

Class Imbalance



Weighted
cross
entropy

Hard Differentiability



Vision reconstitution
on the Retain Set
through 3 epochs

Slight confusion &
a final stabilising
epoch



Model
Unlearned!

→ The confusion process is instant. The whole computation is dedicated to the vision reconstitution (Time Efficient).

3rd place solution - Seif Eddine Achour

Paper approach: Loss landscape adjuster

Regression



The model forgot totally about the unwanted data despite its big size




Original r-squared = 0.28 vs Unlearned r-squared = 0.97!

3rd place solution - Seif Eddine Achour

Paper approach: Loss landscape adjuster

Classification



- Good results without considering the Forget Set (1st approach)



The proper use of the Forget Set will certainly improve results


- The retrained model is not always the ideal one



-1% accuracy on Retain Set lead to -24% of accuracy on Forget Set



The metric which check the similitude between the unlearned and retrained model is not that representative for the unlearning performance



2nd place solution - [kookmin Univ] LD&BGW&KJH

Gradient-based re-initialization method

We assumed that if the gradients of the weights in the model, specifically in the retain set and forget set, are similar, it becomes challenging to forget information from the forget set during the retraining of the retain set.

proposed gradient-based re-initialization method for unlearning consists of three main steps:

1. Gradient Collection:

Gradient information is collected from the forget set and the retain set.

2. Weight initialization:


Based on the gradient information collected in the first step, a percentage of the convolution filter weights are re-initialized.

3. Retraining:

The model is retrained with the retain set. The learning rate for the Uninitialized weights uses 1/10 of the base learning rates.


2nd place solution - [kookmin Univ] LD&BGW&KJH

1. Gradient Collection



- Collect gradients of forget set using gradient ascent
- Collect gradients of retain set using gradient descent
- Random sampling was used from the retain set to match the number of samples in the forget set
- In short, this is simply subtracting forget set's gradient from the retain set's gradient


2. Weight initialization



- Based on the gradient information a percentage of the convolution filter weights are re-initialized
- Our best method re-initialized 30% of convolution filter weights

2nd place solution - [kookmin Univ] LD&BGW&KJH

3. Retraining



- Re-initialized Model is trained using the retain set
- Learning rate for the uninitialized weights uses 1/10 of the base learning rates (accomplished by scaling the gradient of uninitialized weight)
- Used a linear decay learning rate scheduler with a few warmup epoch
 - Consistently produces better results than other learning rate schedulers
 - Used warmup epoch of 3
(0.00033 to 0.001 in the first 3 epochs, and then linearly decreases from 0.001 to 0.00033 in the last 2 epochs)

1st place solution - fanchuan

1. **Forget phase**: minimize KL-divergence between output logits and a uniform pseudo label on forget set.
1. **Adversarial fine-tuning phase**. Alternate between "forget" and "retain" rounds:

Forget round: Maximize dissimilarity between logits of forget and retain set

$$l_i = -\frac{1}{batchsize_2} \sum_{t=0}^{batch_2} \log\left(\frac{e^{sim(x_i, y_t)/\tau}}{\sum_{j=0}^{batch_2} e^{sim(x_i, y_j)/\tau}}\right)$$
$$L_{forget} = \frac{1}{batchsize_1} \sum_{i=0}^{batch_1} l_i$$

Retain: original loss (cross entropy) on retain set.

Trick: Increase batch size from 64 to 258 to be able to perform more epochs (6 -> 8)